

apidays  
**10**  
YEARS  
ANNIVERSARY

apidays  
NEW YORK

API-driven Regulations for Finance,  
Insurance, and Healthcare

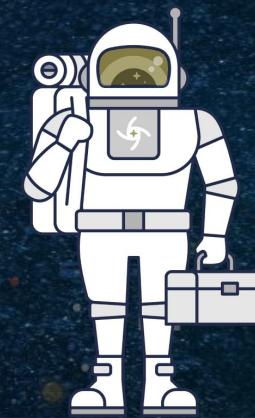
July 27 & 28, 2022

# Expose Rest, Graphql, Grpc Apis on Top for Your Databases

DataStax  
Developers



STARGATE



# Cédrick Lunven

## Director Developer Relations



@clunven



@clun



SDK



- Trainer
- Public Speaker
- Developers Support
- Developer Applications
- Developer Tooling



- Creator of ff4j (ff4j.org)
- Maintainer for 8 years+



{ REST }

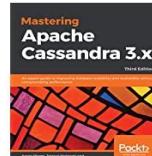
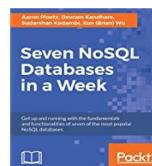


gRPC

- Happy developer for 14 years
- Spring Petclinic Reactive & Starters
- Implementing APIs for 8 years

# Aaron Ploetz

DBRE/Developer Advocate



@aaronploetz



@aploetz



@aploetz

- DB Engineer, Author
- Former SWE/DevOps/DB Lead @ W.W. Grainger & Target
- Rep leader of S.O. "Cassandra" tag
- Worked as an author on:
  - Mastering Apache Cassandra 3.x
  - Seven NoSQL Databases in a Week

# Materials for today

[https://github.com/  
datastaxdevs/  
conference-  
Expose-Rest-Graphql-Grpc-Apis-for-Databases](https://github.com/datastaxdevs/conference-Expose-Rest-Graphql-Grpc-Apis-for-Databases)



# Agenda

**01**

**Data Api**  
**What, Why**

**02**

**Rest, GraphQL, gRPC**  
**SWOT**

**03**

**Data Gateway**  
**Why and How**

**04**

**Stargate.io**  
**Architecture**

**05**

**Demo**  
**Api and Tooling**

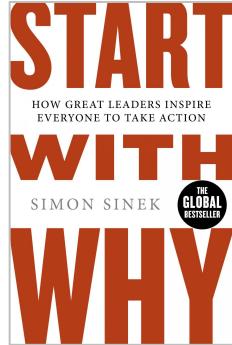
**05**

**What's NEXT ?**  
**Resources**

# Data API : Start with Why

- Developers

- Do you like learning query languages (CQL, N1QL, GQL, cypher, gremlin....)
  - No. Save my JSON, give it to me back when I need it
  - ORM / Spring Data are so popular nowadays
- Do you care about how your data is stored ?
  - Physical data model part of the interface...yikes
  - Create structures based on queries
- Do you like installing and running databases locally
  - especially distributed databases
  - especially with datasets and integration tests



# Data API : Start with Why

- Operators and Databases Administrators...

- Do you allow developers to execute direct queries against your DB ?
- Do you like opening port ranges like 0-65536 to allow communications with applications, especially in the cloud.
- Do you like creating dedicated projects and hiring people just to create APIs to expose an existing treatment (digital transformation FTW)
- Security, Monitoring, Throttling



Me when devs ask for admin access to the database.

# Agenda

01

Api and Databases  
What, Why and How

02

Rest, GraphQL, gRPC  
SWOT

03

Data Gateway  
Why and How

04

Stargate.io  
Architecture

05

Demo  
Api and Tooling

05

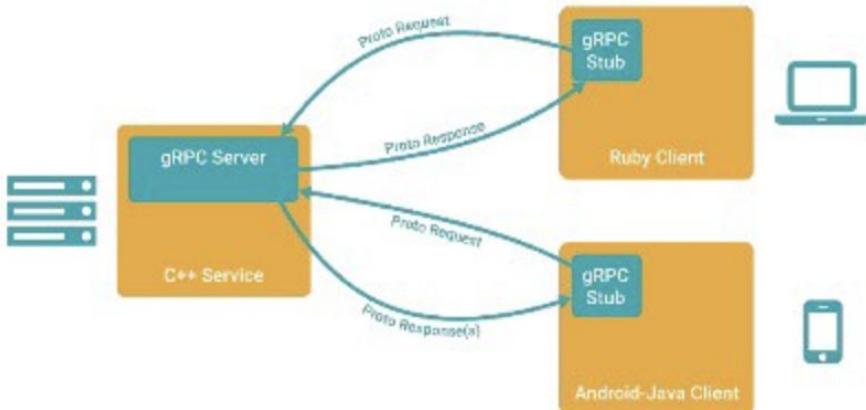
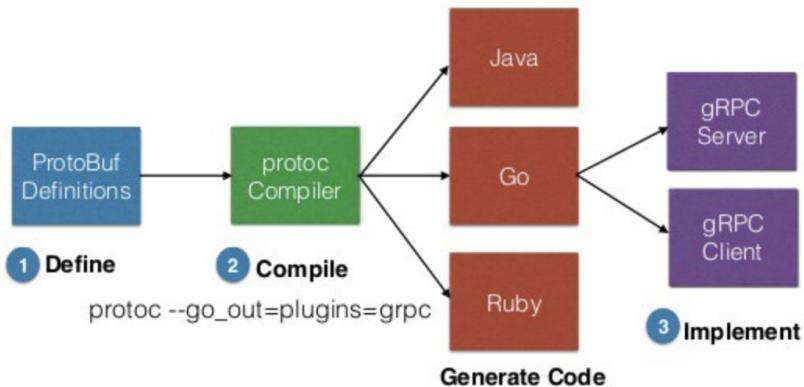
What's NEXT ?  
Resources

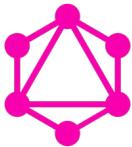
# { REST }

The screenshot shows the Swagger UI interface for a REST API. At the top, there's a logo for "Swagger Supported by SMARTBEAR" and a link to "/swagger.json". On the right, there's a green "Explore" button and a dropdown arrow. The main area is titled "documents". Below the title, a list of API endpoints is displayed in colored boxes:

- GET** /v2/namespaces/{namespace-id}/collections List collections in namespace
- POST** /v2/namespaces/{namespace-id}/collections Create a new empty collection in a namespace
- GET** /v2/namespaces/{namespace-id}/collections/{collection-id} Search documents in a collection
- POST** /v2/namespaces/{namespace-id}/collections/{collection-id} Create a new document
- DELETE** /v2/namespaces/{namespace-id}/collections/{collection-id} Delete a collection in a namespace
- POST** /v2/namespaces/{namespace-id}/collections/{collection-id}/upgrade Upgrade a collection in a namespace
- POST** /v2/namespaces/{namespace-id}/collections/{collection-id}/batch Write multiple documents in one request
- GET** /v2/namespaces/{namespace-id}/collections/{collection-id}/{document-id} Get a document
- PUT** /v2/namespaces/{namespace-id}/collections/{collection-id}/{document-id} Create or update a document with the provided document-id
- DELETE** /v2/namespaces/{namespace-id}/collections/{collection-id}/{document-id} Delete a document
- PATCH** /v2/namespaces/{namespace-id}/collections/{collection-id}/{document-id} Update data at the root of a document
- GET** /v2/namespaces/{namespace-id}/collections/{collection-id}/{document-id}/{document-path} Get a path in a document
- PUT** /v2/namespaces/{namespace-id}/collections/{collection-id}/{document-id}/{document-path} Replace data at a path in a document
- DELETE** /v2/namespaces/{namespace-id}/collections/{collection-id}/{document-id}/{document-path} Delete a path in a document
- PATCH** /v2/namespaces/{namespace-id}/collections/{collection-id}/{document-id}/{document-path} Update data at a path in a document
- GET** /v2/namespaces/{namespace-id}/collections/{collection-id}/json-schema Get a JSON schema from a collection

# gRPC





# GraphQL

It is a web protocol (HTTP) that specifies the way we build and query remote APIs using a **Tree/JSON Syntax**.

Based on a **strongly typed** language named GraphQL SDL (Schema Definition Language)

```
incomingWorkshops {  
  title  
  abstract  
  speakers  
}
```

```
type Workshop {  
  title: String!  
  abstract: String  
  speakers: [Speaker]  
  releaseYear: int  
}
```



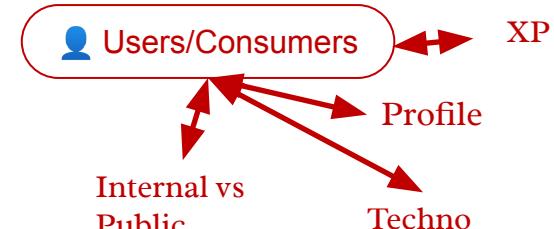
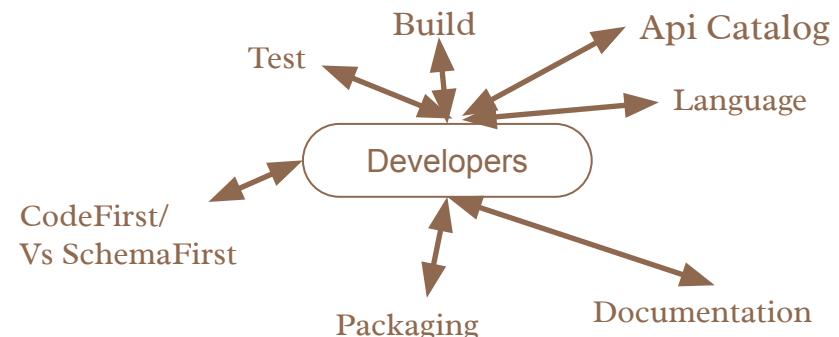
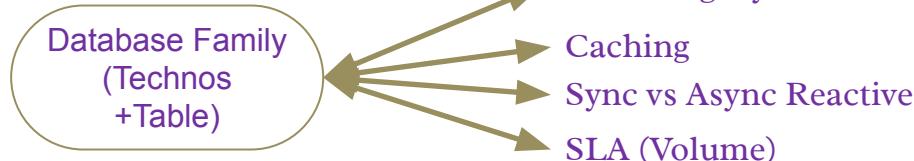
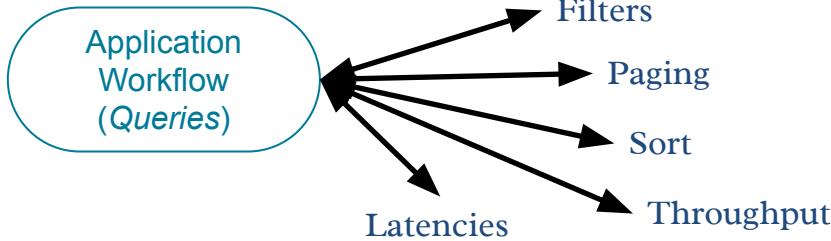
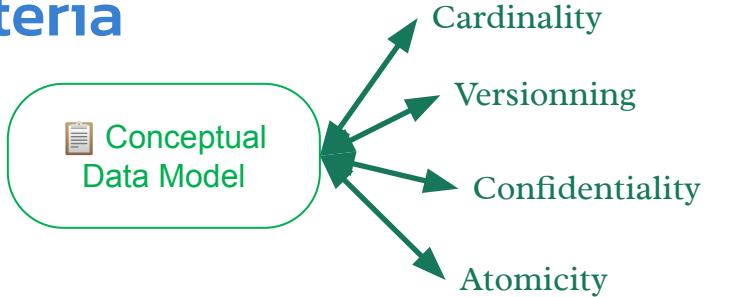
A protocol that allows the client to specify exactly **what** data it needs from a model

```
incomingWorkshops {  
    title  
}
```

It allows one to aggregate data from **multiple relations** in a single query

```
incomingWorkshops {  
    title: String!  
    abstract: String  
    speakers: {  
        name  
    }  
    releaseYear: int  
}
```

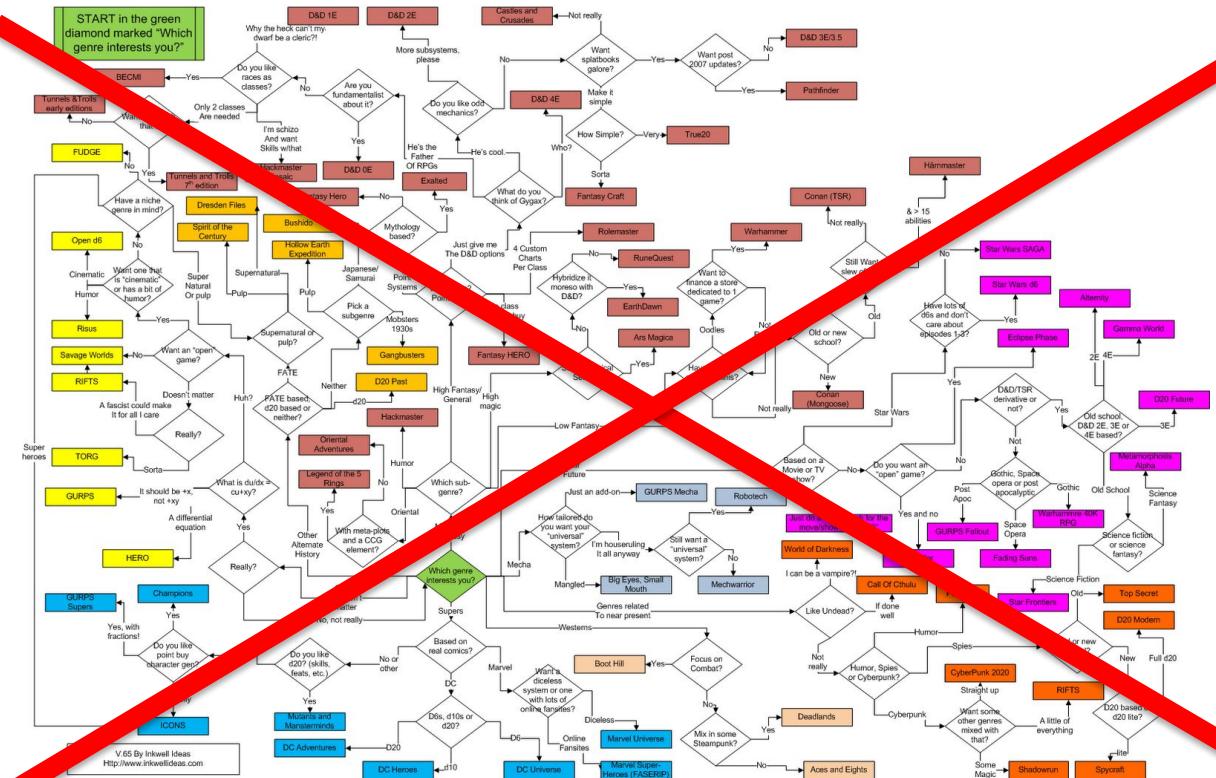
# Criteria



# Matrix Analysis

Concept	GraphQL	REST	gRPC
<b>Modèle conceptuel de données (Entités, Relations)</b>			
Cardinalité	1	5	3
Versionning	1	8	9
Confidentialité	6	9	1
Atomicité	6	7	0
<b>Parcours Applicatifs</b>			
Filtres	2	3	4
Pagination	3	3	6
Tris	4	4	6
Throughput	4	3	6
Latence	6	4	6
Base de données	2	3	6
Intégrité	6	4	7
Cache	7	0	6
Sync vs Async vs Reactive	1	2	4
SLA	6	3	4
<b>Development</b>			
Language	1	4	5
Api Catalog	6	7	8
Test	3	2	0
Build	4	5	6
Packaging	4	6	7
Code first / Schema Frist	5	7	8
Documentation	2	5	1
<b>Client</b>			
XP	1	4	7
Profil	3	3	9
Techno	4	5	7
Interne vs Extern	4	5	8

# Decision Tree



# { REST }



- ❖ Decoupling Client / Server (*Schema on read*)
- ❖ Api Lifecycle (*Versioning*)
- ❖ Tooling (*API Management, Serverless*)



- ❖ Verbose payloads (*json, xml*)
- ❖ No discoverability
- ❖ Not suitable for command-like (functions) API



- ❖ CRUD superstar
- ❖ Relevant for mutations (OLTP)
- ❖ Public and web APIs
- ❖ Limited Business Scope



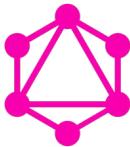
- ❖ High Performances (*http/2 – binary serialisation*)
- ❖ Multiple stubs : Sync, Async, Streaming
- ❖ Multi languages - Interoperability



- ❖ Strongly coupled (*schema with proto files*)
- ❖ No discoverability
- ❖ *Protobuf* serialization format



- ❖ Distributed network of services (no waits)
- ❖ High throughput & streaming use cases
- ❖ Command-like (eg: *slack*)



# GraphQL



- ❖ Discoverability, documentation
- ❖ Custom payloads
- ❖ Match standards (Json | Http)



- ❖ Single endpoint (*versioning, monitoring, security*)
- ❖ Complex implementation (*tooling, still young*)
- ❖ Nice for customers nasty for DB (*N+1 select*)



- ❖ Backend for frontend (JS)
- ❖ Service aggregation | composition (*joins*)
- ❖ When volume matters (*mobile phones*)

# Agenda

01

Api and Databases  
What, Why and How

02

Rest, GraphQL, gRPC  
SWOT

03

Data Gateway  
Why and How

04

Stargate.io  
Architecture

05

Demo  
Api and Tooling

05

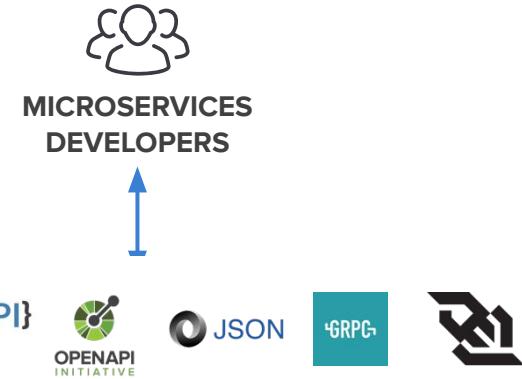
What's NEXT ?  
Resources

# Data Gateway



Developers want the option to use modern APIs and development gateways.

Apache Cassandra is a distributed database designed for the new standard and the associated APIs that facilitate the first choice of developers.



DataStax Developers

# Stargate Overview

An open source API framework for data

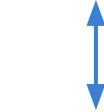
Stargate makes it easy to use a database for any application workload by adding plugin support for new APIs, data types, and access methods



MICROSERVICES  
DEVELOPERS



STARGATE



cassandra

# Stargate.io

stargate / stargate Public

Code Issues 246 Pull requests 13 Discussions Actions Projects Wiki Security ...

# OPEN SOURCE DATA GATEWAY

Stargate is an open source data gateway that sits between your app and your databases. Stargate brings together an API platform and data request coordination code into one OSS project.

[Get Started](#)  [Github](#)

Stargate is the official data API for DataStax Astra and DataStax Enterprise! [Try It Now](#)

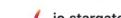
Reference example for gRPC java client (#1236) [View file](#) [Edit file](#) [Delete file](#) [Revert file](#) [Compare file](#) [Pull Request](#) [Merge pull request](#) [Create merge commit](#) [Create a new branch](#) [Create a new pull request](#) [Create a new issue](#) [Create a new discussion](#) [Create a new project](#) [Create a new wiki page](#) [Create a new security issue](#)

About  
An open source data gateway

[Readme](#) [Apache-2.0 License](#)

Releases 62  
 [Release v1.0.32](#) [Latest](#) [25 days ago](#) [+ 61 releases](#)

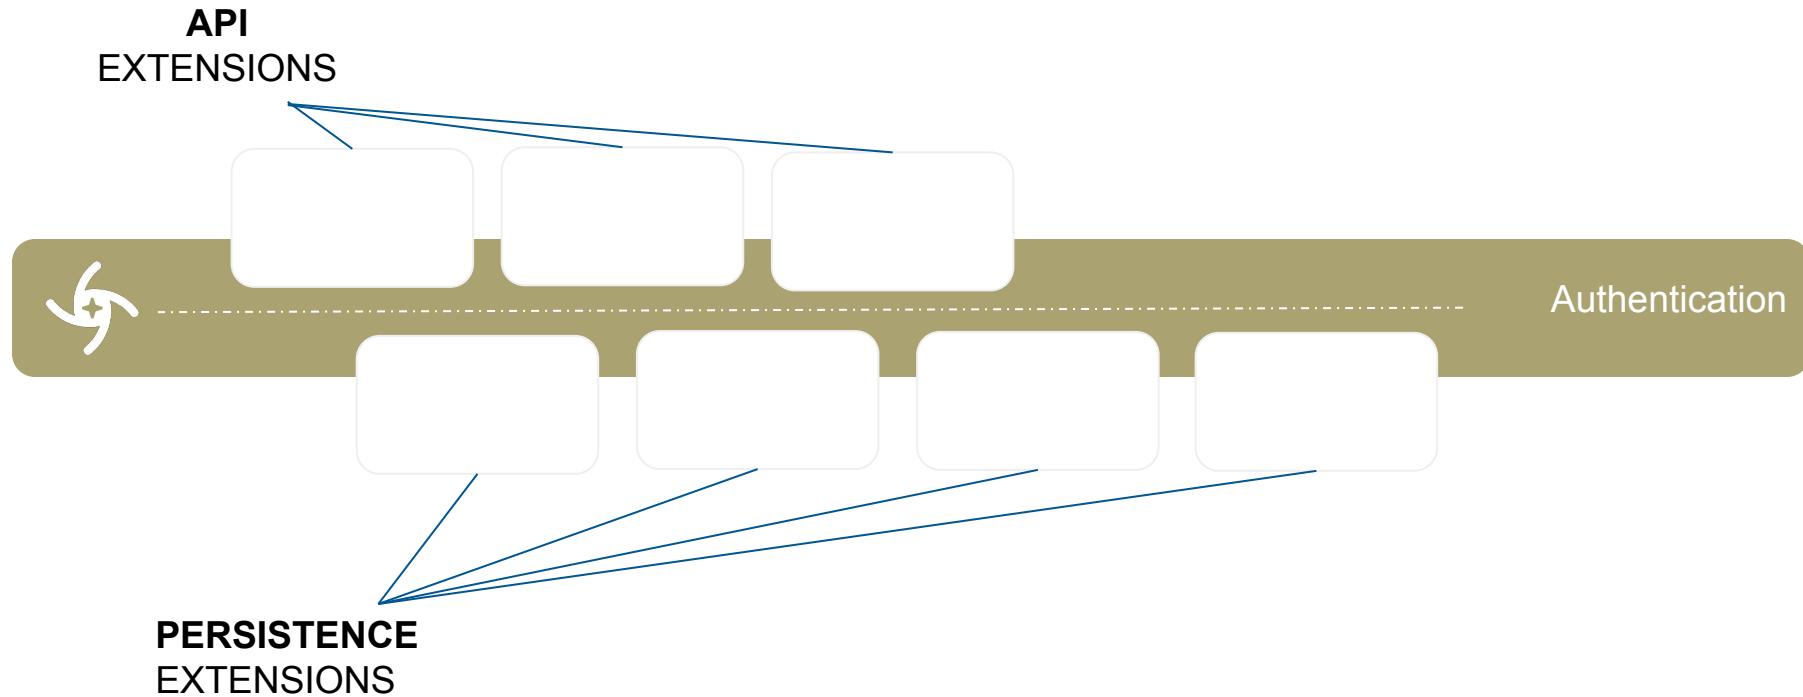
Packages 13  
   [+ 10 packages](#)

Contributors 30 

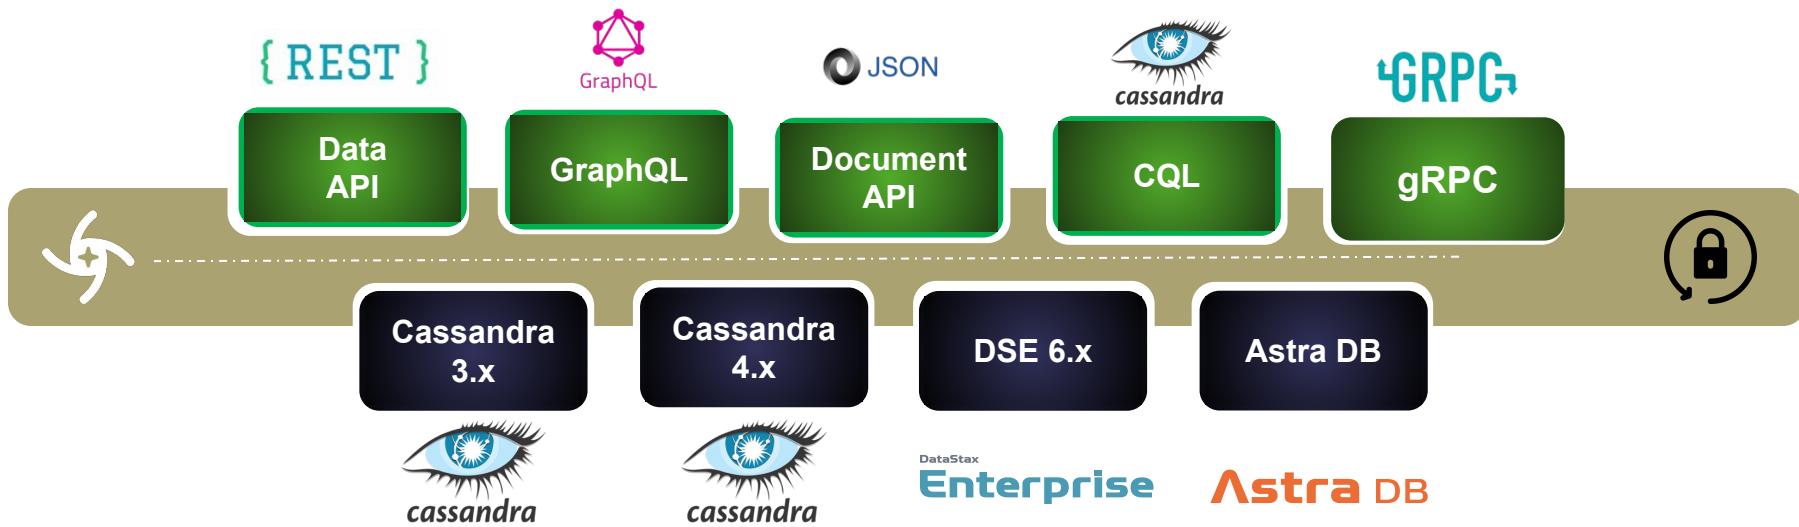
apidays

DataStax Developers

# API and Persistence Extensions



# API Extensions and Persistence Extensions



# Agenda

**01**

**Api and Databases**  
**What, Why and How**

**02**

**Rest, GraphQL, gRPC**  
**SWOT**

**03**

**Data Gateway**  
**Why and How**

**04**

**Stargate.io**  
**Architecture**

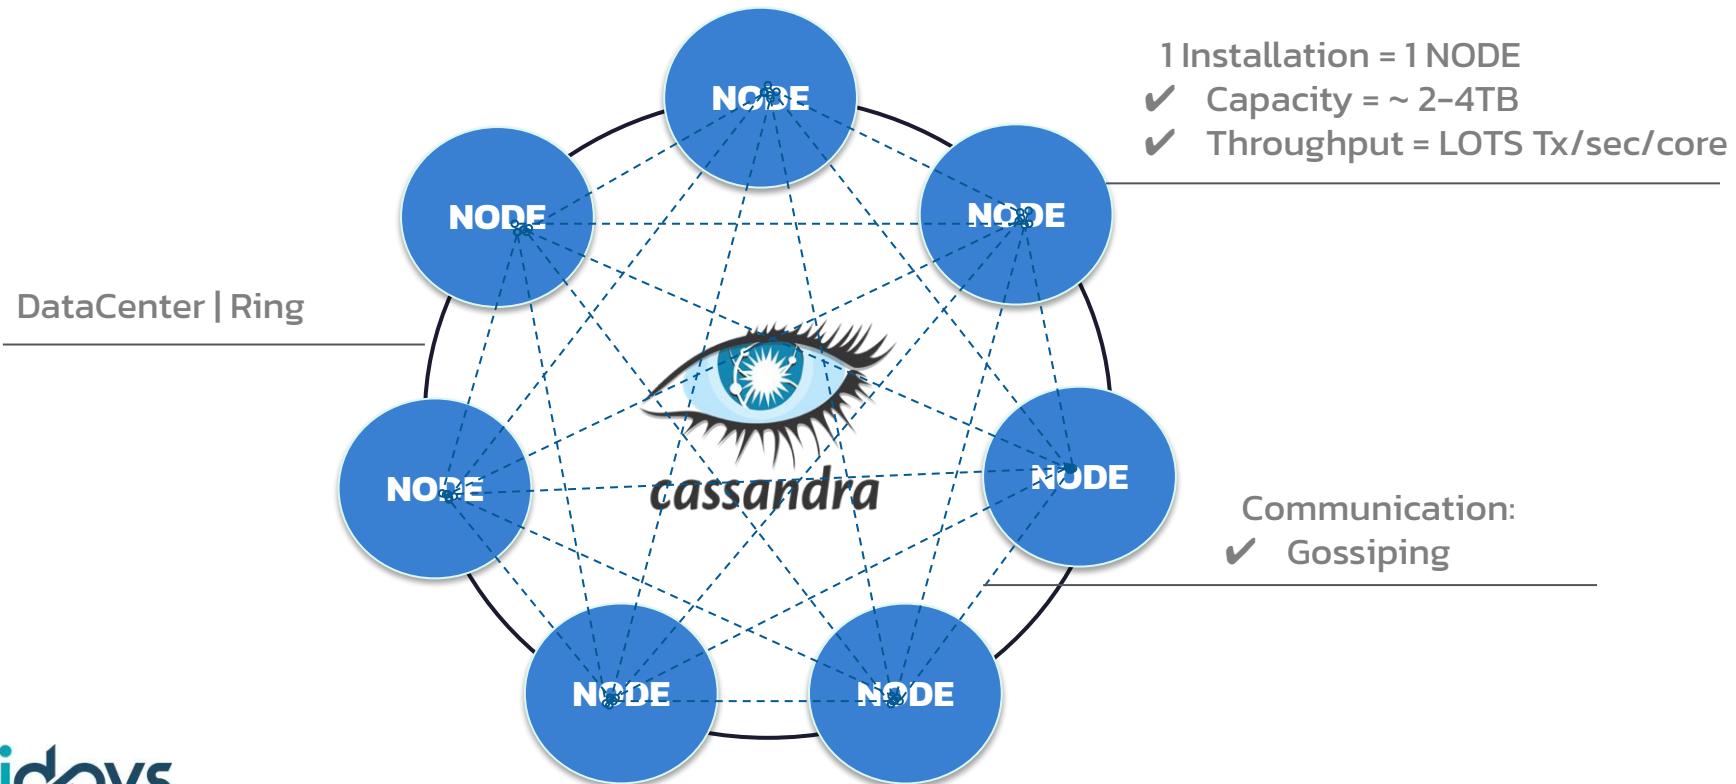
**05**

**Demo**  
**Api and Tooling**

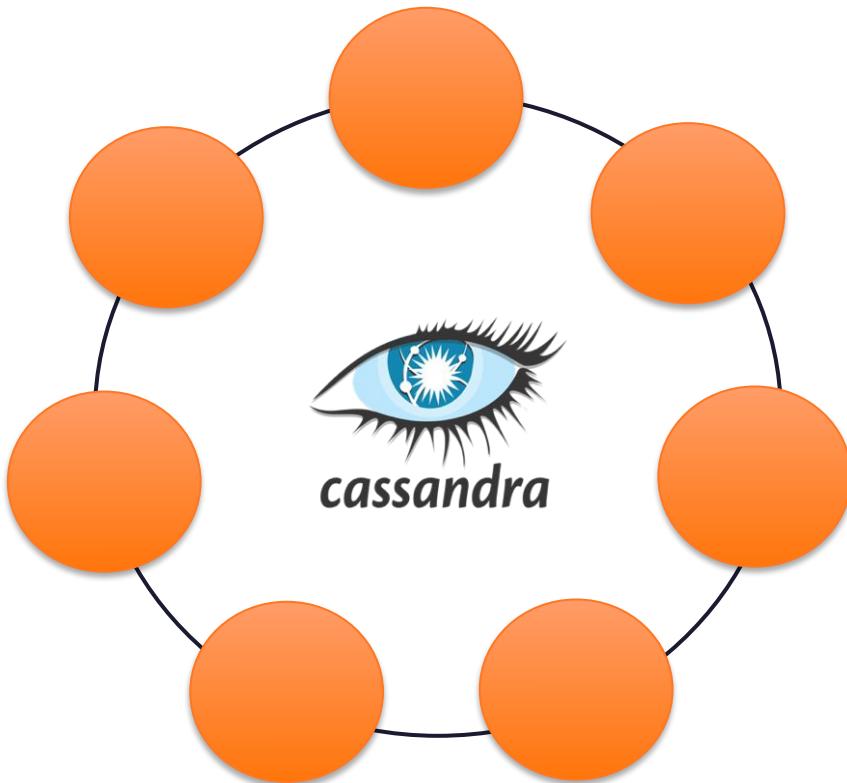
**05**

**What's NEXT ?**  
**Resources**

# Apache Cassandra™ = NoSQL Distributed Database



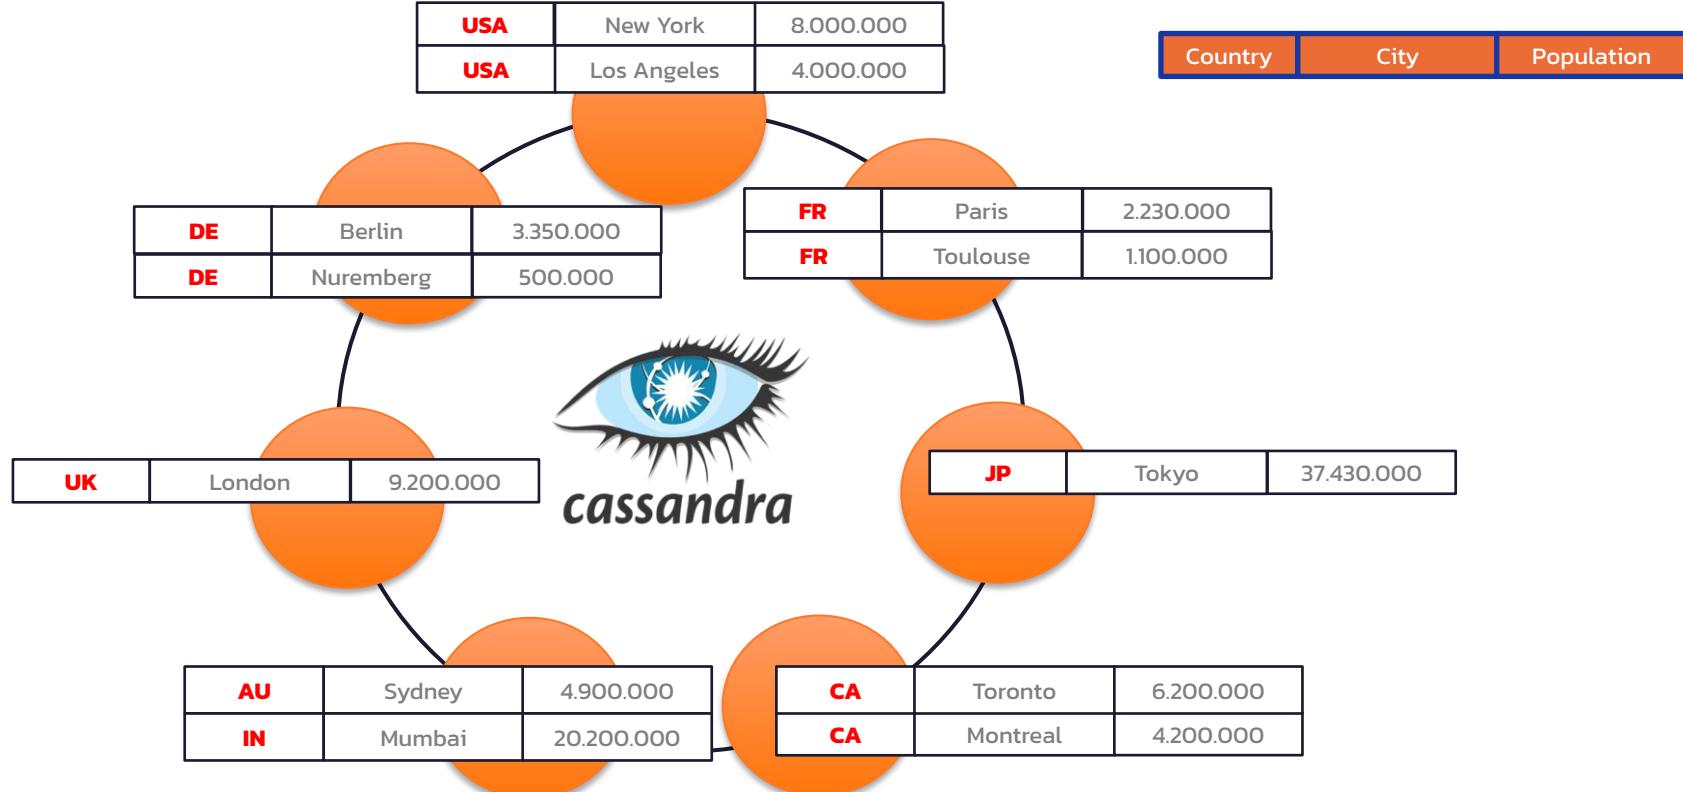
# Data is Distributed



Country	City	Population
USA	New York	8.000.000
USA	Los Angeles	4.000.000
FR	Paris	2.230.000
DE	Berlin	3.350.000
UK	London	9.200.000
AU	Sydney	4.900.000
DE	Nuremberg	500.000
CA	Toronto	6.200.000
CA	Montreal	4.200.000
FR	Toulouse	1.100.000
JP	Tokyo	37.430.000
IN	Mumbai	20.200.000

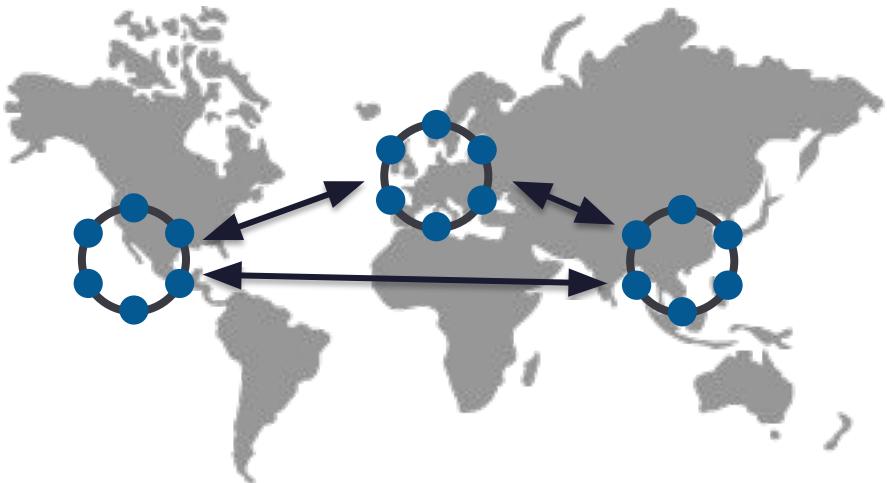
Partition Key

# Data is Distributed

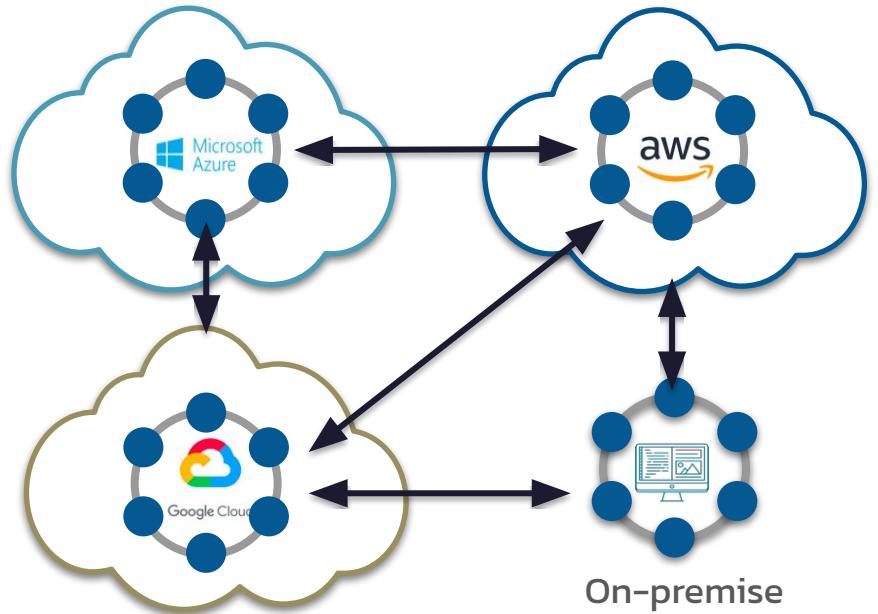


# Data Distributed Everywhere

- Geographic Distribution



- Hybrid-Cloud and Multi-Cloud



# Understanding Use Cases

Scalability



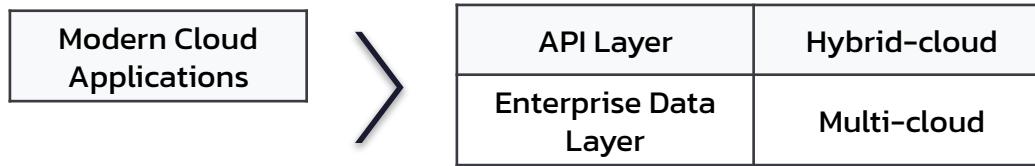
Availability



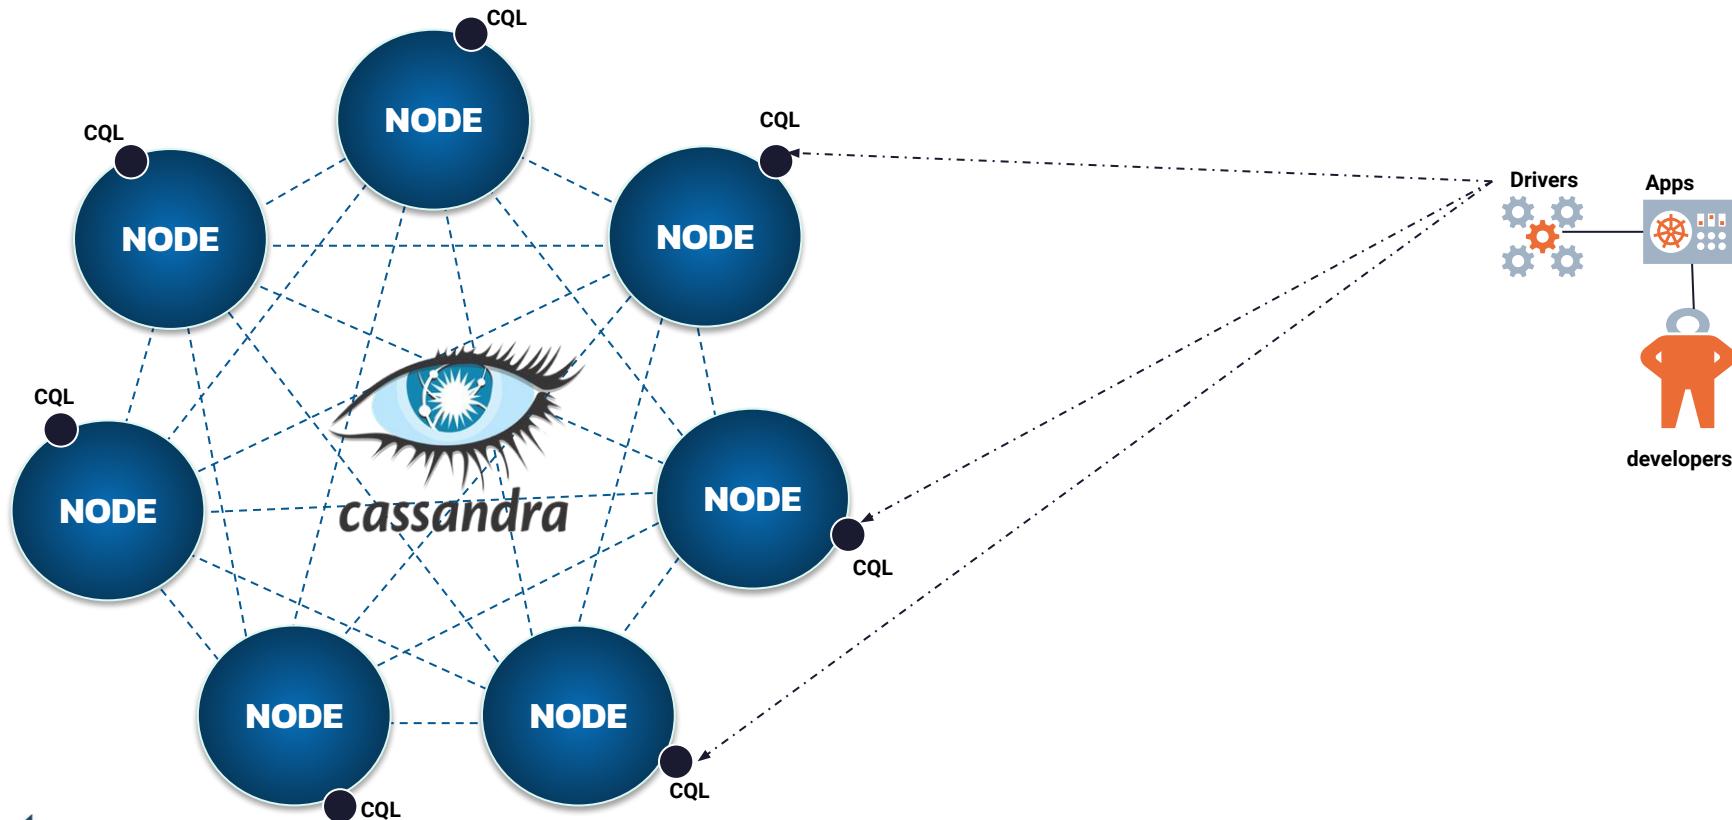
Distributed



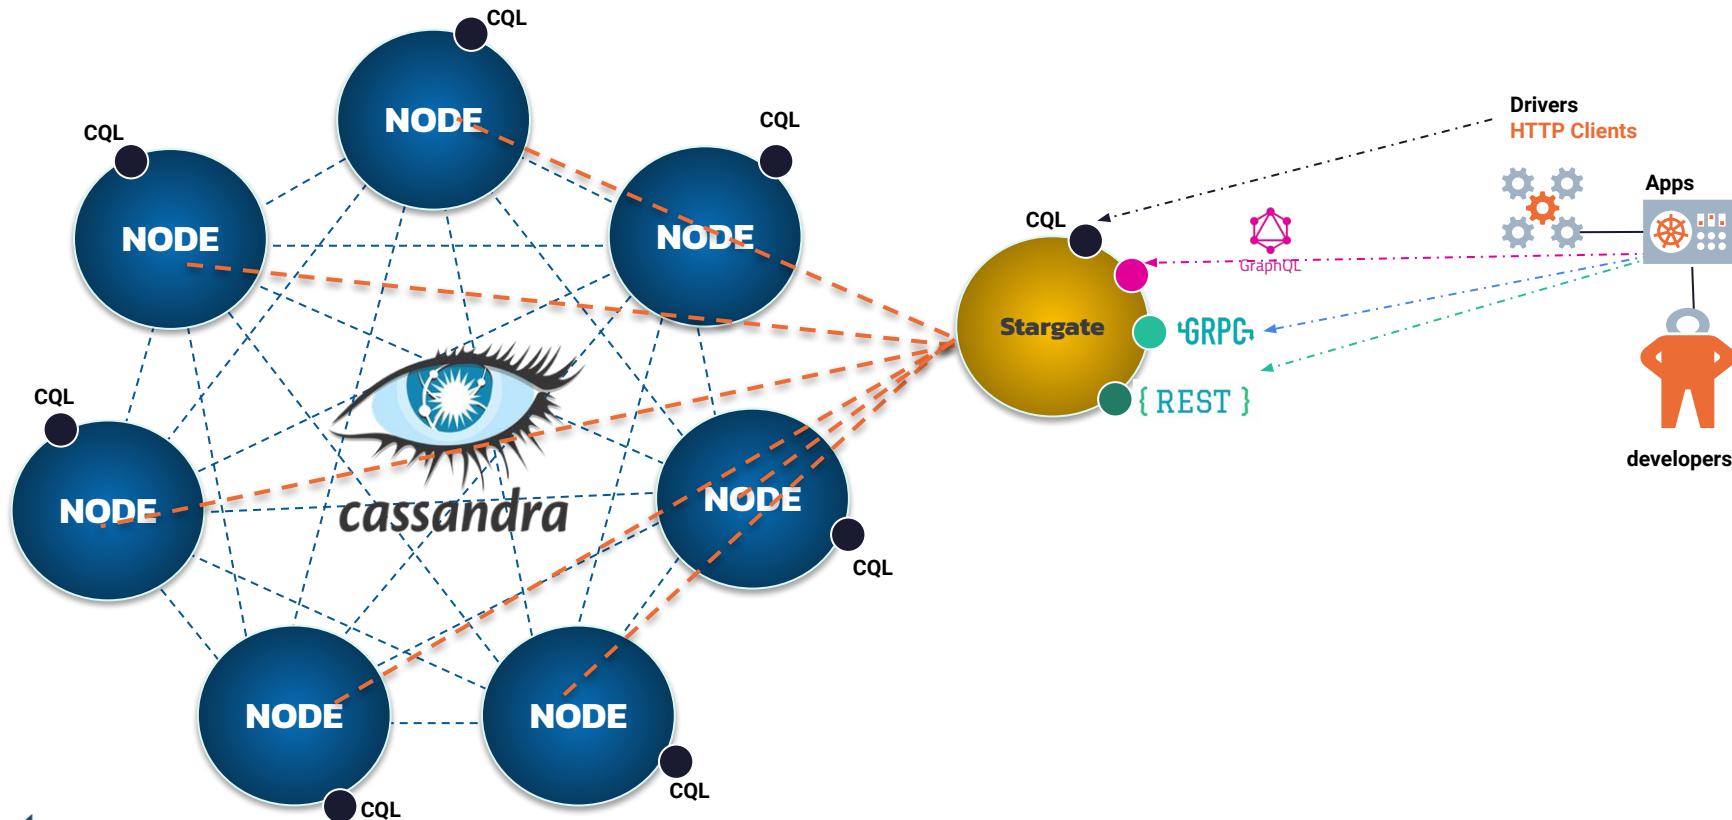
Cloud-native



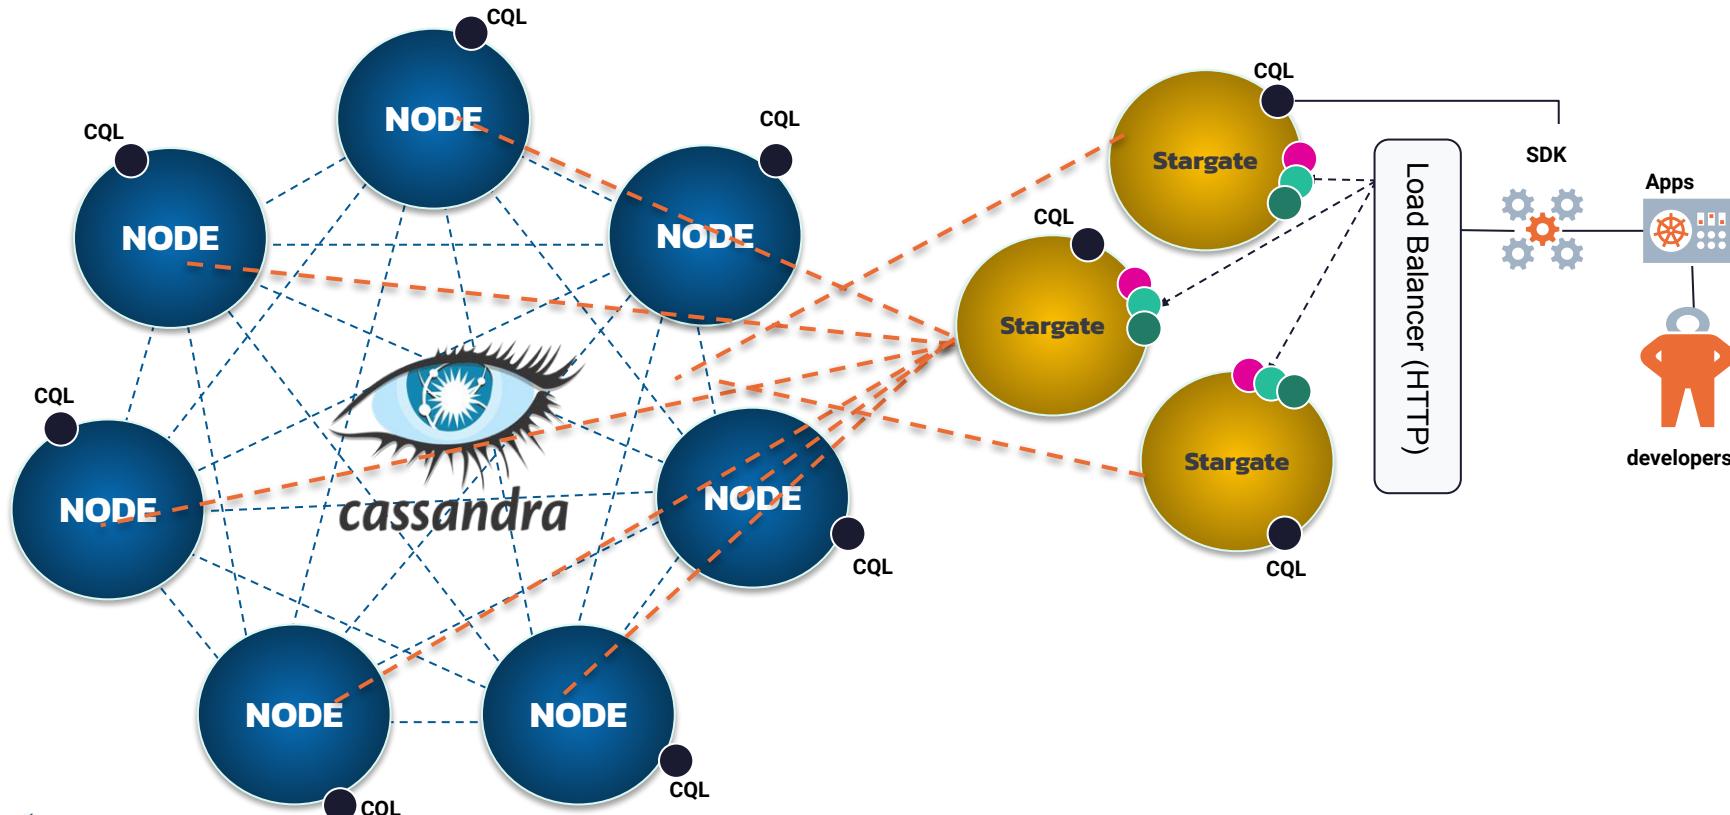
# Connecting to your cluster (Before)



# Connecting to your cluster (with Stargate)



# Connecting to your cluster (with Stargates)



# Agenda

**01**

**Api and Databases**  
**What, Why and How**

**02**

**Rest, GraphQL, gRPC**  
**SWOT**

**03**

**Data Gateway**  
**Why and How**

**04**

**Stargate.io**  
**Architecture**

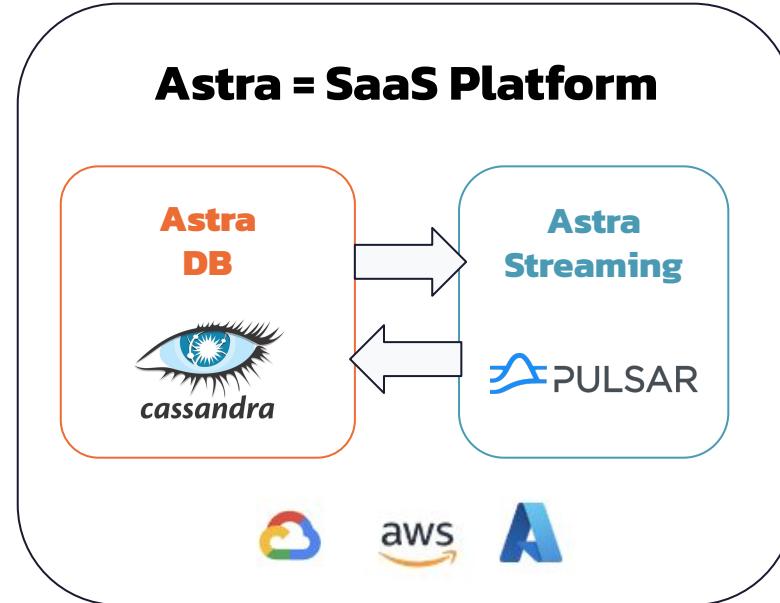
**05**

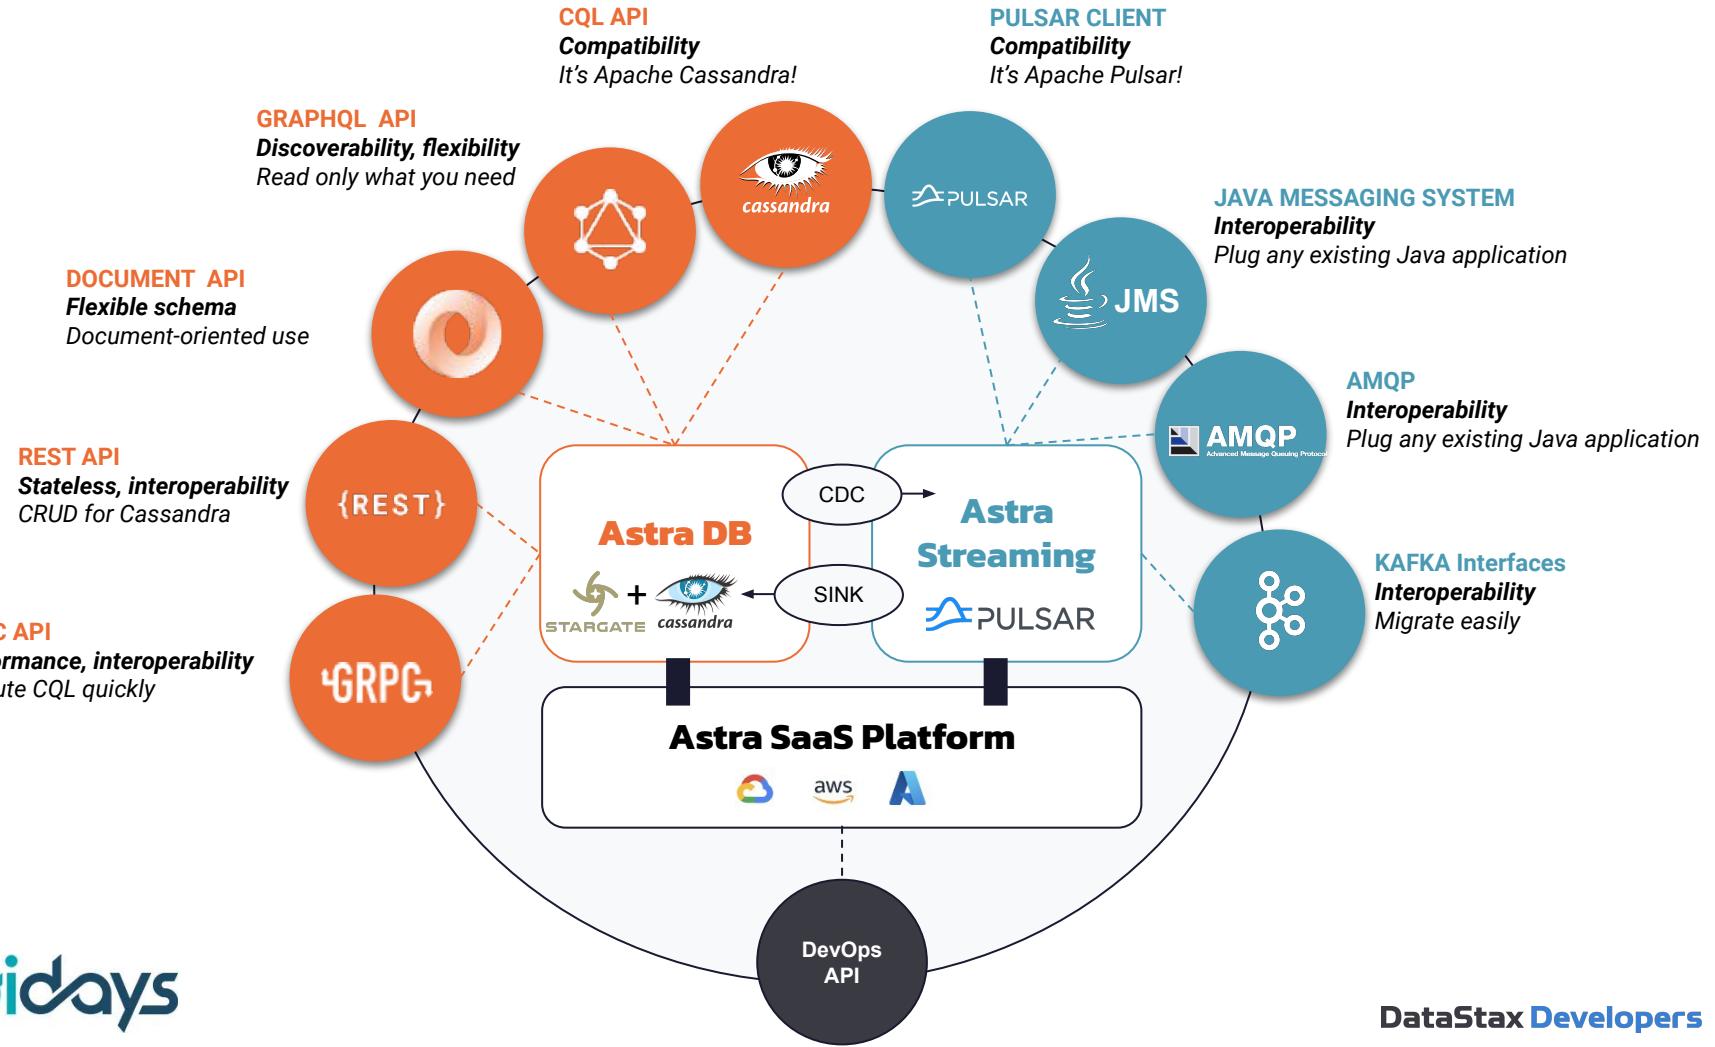
**Demo**  
**Api and Tooling**

**05**

**What's NEXT ?**  
**Resources**

# Astra DB





# HouseKeeping

- Instructions and Slide
  - <https://github.com/datastaxdevs/conference-Expose-Rest-Graphql-Grpc-Apis-for-Databases>
- Runtime
  - <https://astra.datastax.com>



## 1. Create Astra Instance

**ASTRA** is the simplest way to run Cassandra with zero operations at all - just push the button and get your cluster. No credit card required, \$25.00 USD credit every month, roughly 5M writes, 30M reads, 40GB storage monthly - sufficient to run small production workloads.

Register (if needed) and Sign In to Astra <https://astra.datastax.com>: You can use your Github, Google accounts or register with an email.

*Make sure to chose a password with minimum 8 characters, containing upper and lowercase letters, at least one number and special character*

Create a "pay as you go" plan

Follow this [guide](#), to set up a pay as you go database with a free \$25 monthly credit.

- Select the pay as you go option: Includes \$25 monthly credit - no credit card needed to set up.

You will find below which values to enter for each field.

- For the database name - `free_db`. While Astra allows you to fill in these fields with values of your own choosing, please follow our recommendations to ensure the application runs properly.

- For the keyspace name - `free`. It's really important that you use the name "free" for the code to work.

*You can technically use whatever you want and update the code to reflect the keyspace. This is really to get you on a happy path for the first run.*

- For provider and region: Choose and provider (either GCP or AWS). Region is where your database will reside physically (choose one close to you or your users).

# Hands-on

## #1/2 Getting your instance

### 1. Create Astra Instance

Astra is the simplest way to run Cassandra with zero operations at all - just push the button and get your cluster. No credit card required, \$25.00 USD credit every month, roughly 5M writes, 30M reads, 40GB storage monthly - sufficient to run small production workloads.

- Register (if needed) and Sign In to Astra <https://astra.datastax.com>: You can use your Github, Google accounts or register with an email.

*Make sure to choose a password with minimum 8 characters, containing upper and lowercase letters, at least one number and special character*

- Create a "pay as you go" plan

Follow this [guide](#), to set up a pay as you go database with a free \$25 monthly credit.

- **Select the pay as you go option:** Includes \$25 monthly credit - no credit card needed to set up.

You will find below which values to enter for each field.

- **For the database name** - `free_db`. While Astra allows you to fill in these fields with values of your own choosing, please follow our recommendations to ensure the application runs properly.
- **For the keyspace name** - `keyspace1`. It's really important that you use the name "free" for the

### 2. Working with Cassandra

- Check that our keyspace exist

```
describe keyspaces;
```

- Create Entities

```
use keyspace1;
```

```
CREATE TYPE IF NOT EXISTS video_format (
    width   int,
    height  int
);
```

```
CREATE TABLE IF NOT EXISTS videos (
    videoid  uuid,
    title    text,
    upload   timestamp,
    email    text,
    url      text,
    tags     set <text>,
    frames   list<int>,
    formats  map <text,frozen<video_format>>,
    PRIMARY KEY (videoid)
);
```

```
describe keyspace1;
```

# Agenda

**01**

**Api and Databases**  
**What, Why and How**

**02**

**Rest, GraphQL, gRPC**  
**SWOT**

**03**

**Data Gateway**  
**Why and How**

**04**

**Stargate.io**  
**Architecture**

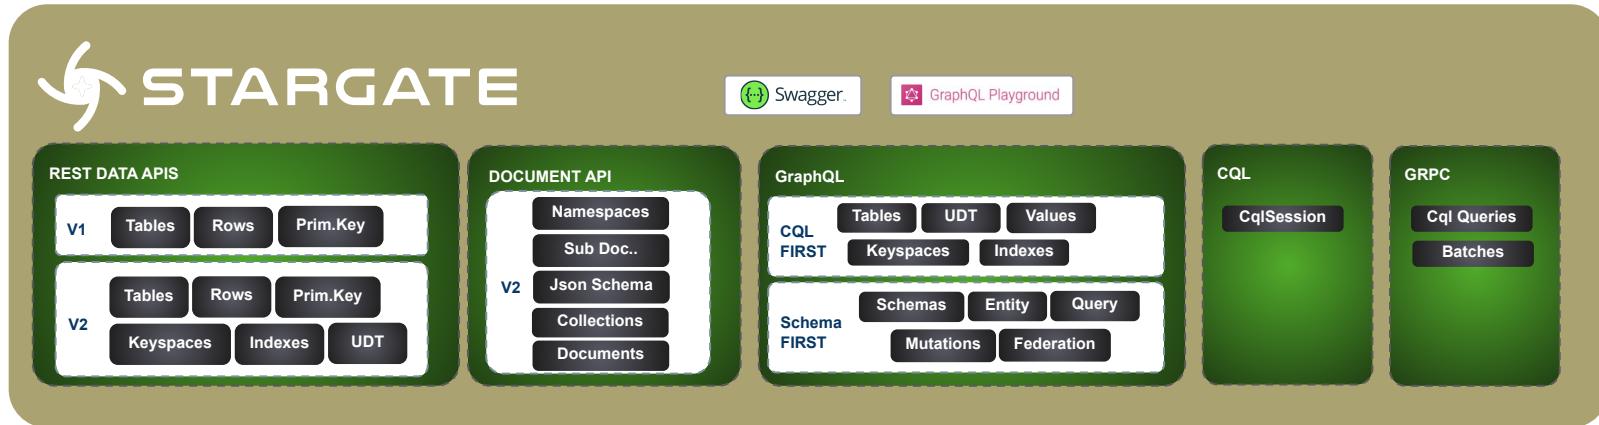
**05**

**Demo**  
**Api and Tooling**

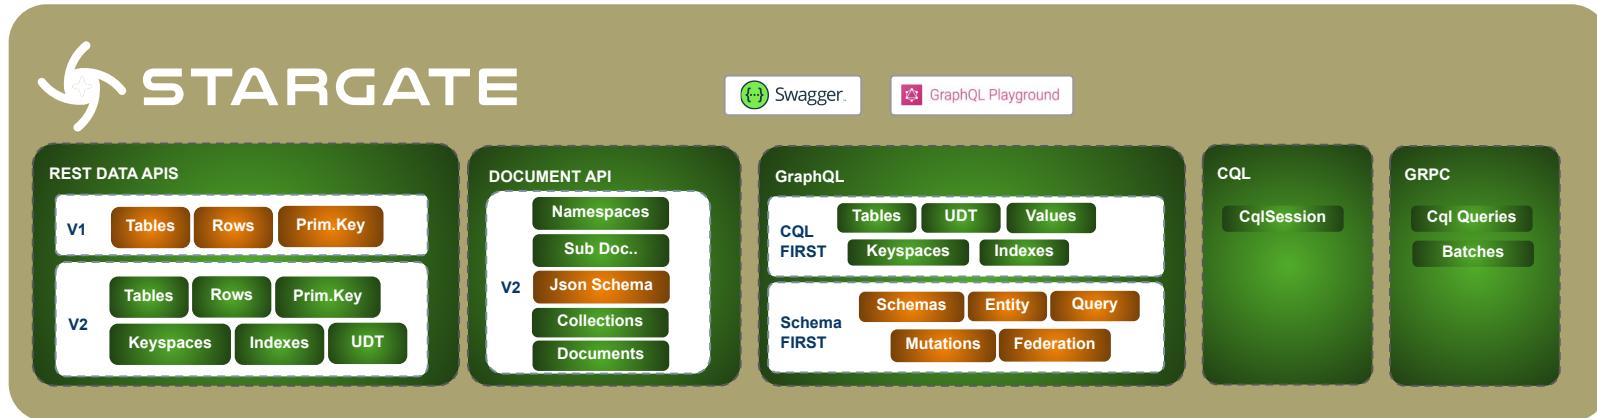
**05**

**What's NEXT ?**  
**Resources**

# Features Map



# SDK in Java, Javascript, Python



# Sample Code

```
StargateClient client = StargateClient.builder()
    .username("k8ssandra-superuser")          // Mandatory username
    .password("JxzrPOnvDGqfEOQ0EySQ")         // Mandatory password
    .endPointAuth("http://localhost:8081")      // Mandatory authentication url, defaulting to http://localhost
    .endPointRest("http://localhost:8082")       // Rest and Document APIs
    .endPointGraphQL("http://localhost:8080")    // GraphQL API

    // Cqlsession Only
    .addCqlContactPoint("127.0.0.", 9042)      // Contact Point
    .localDc("dc1")                            // Local Datacenter is mandatory driver 4xx+
    .keypace("ks1")                            // (optional) Set your keyspace
    .build();
```



# Sample Code

```
// Retrieve an object and marshall
Optional<Address> address = colPersonClient
    .document("e8c5021b-2c91-4015-aec6-14a16e449818")
    .findSubDocument("address", Address.class);

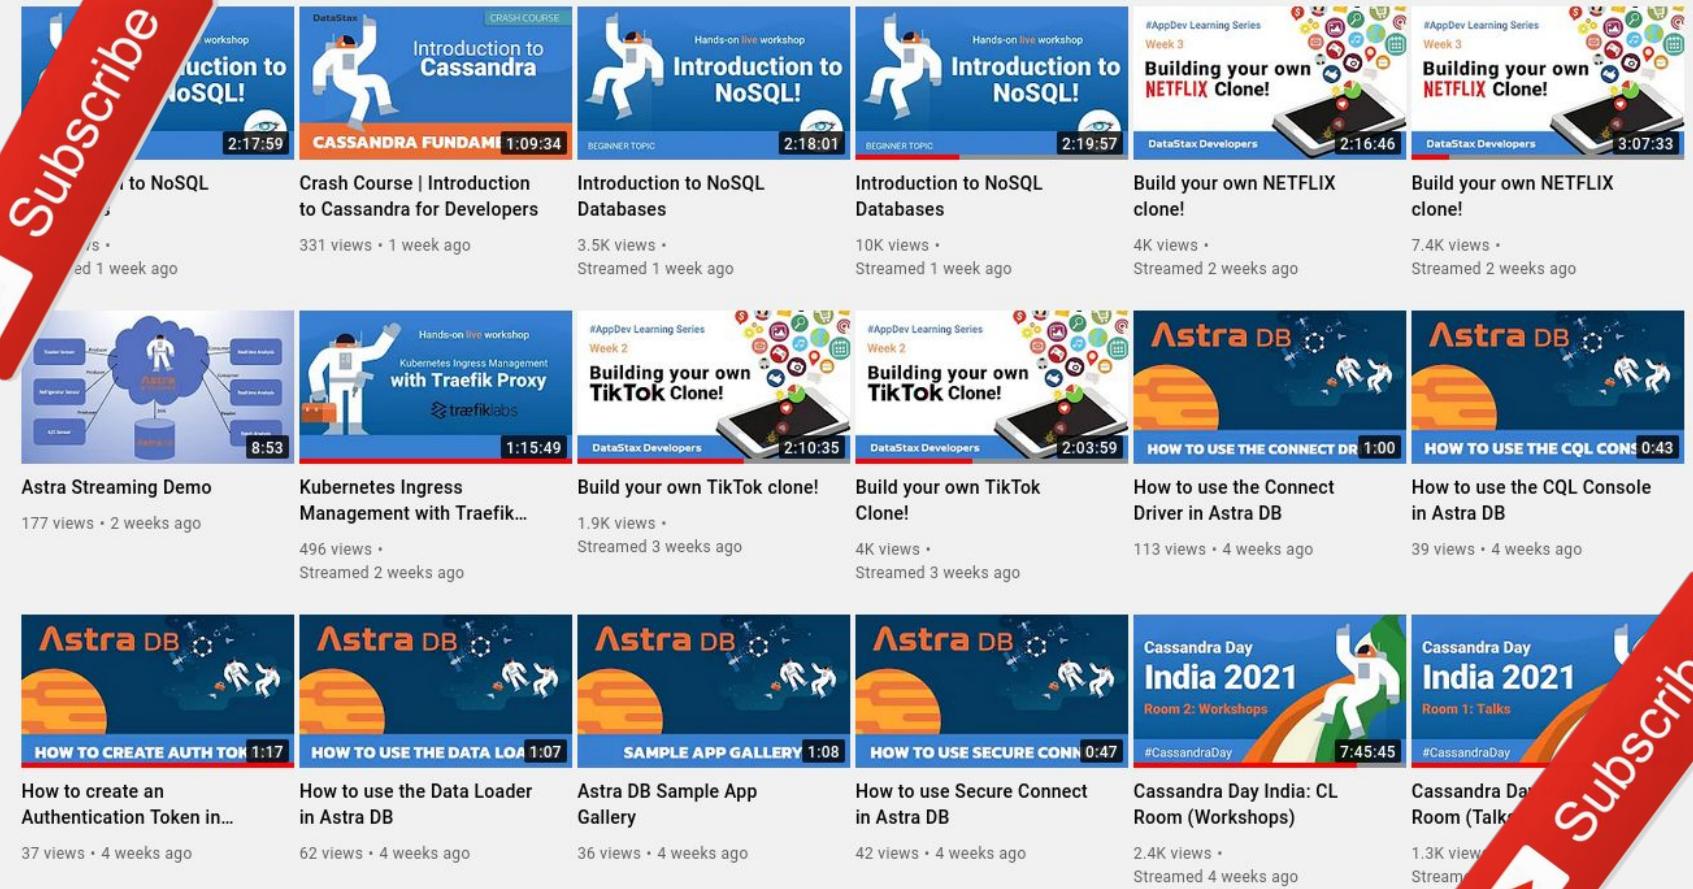
// Building query {"age": {"$gte":30}, "lastname": {"$eq": "PersonAstra2"}}
SearchDocumentQuery query = SearchDocumentQuery.builder()
    .where("age").isGreaterOrEqualsThan(30)      // First filter to use where()
    .and("lastname").isEqualTo("PersonAstra2")   // Any extra filter to use and()
    .withPageSize(10)                            // Default and max pageSize are 20
    .build();

// Retrieve PAGE 1
DocumentResultPage<Person> currentPage = colPersonClient.findPage(query, Person.class);

// Retrieve PAGE 2 (if any)
if (currentPage.getPageState().isPresent()) {
    query.setPageState(currentPage.getPageState().get());
}
DocumentResultPage<Person> nextPage = colPersonClient.findPage(query, Person.class);

// Retrieve all documents in one call (with warning in RED ABOVE)
Stream<Person> allPersons = colPersonClient.findAll(query, Person.class)
```

# Subscribe



# Subscribe

DataStax Developers

# DataStax Developers

## Thank you!



@aploetz



aaronploetz



DataStax Developers